

# Operations ... SQL Specifics



pm @ daiict



# Certain SQL specific features

- Functions for computed attributes in projection
- Certain issues related to NULLs,
  - UNKNOWN is third TRUTH value while evaluating SQL expressions
  - IS NULL as additional predicate
- **UNION/INTERSECT/EXCEPT ALL** in SQL
- Correlated Sub-Queries
- EXISTS
- Some (relation as BAG) comparison operators



# Functions and Operators in SQL



# Functions and Operators

- SQL provides various functions and operators that can be used to create a new attribute in resultant relations
- There are typically, type conversion, arithmetic operators, mathematical, and string manipulation operators and functions. For example: `substring`, `upper`, `lower`, `sqrt`, `ln`, etc.
- Details for PostgreSQL functions can be seen at:  
[http://intranet.daiict.ac.in/~pm\\_jat/postgres/html/functions.html](http://intranet.daiict.ac.in/~pm_jat/postgres/html/functions.html).



# Examples

```
SELECT ssn,  
       fname || ' ' || minit || '. ' || lname AS name,  
       current_date-bdate AS age FROM employee;
```

```
SELECT essn, hours*50 AS amount FROM works_on;
```

```
SELECT upper(fname) AS name, ln(salary) AS x FROM  
employee;
```



# Example

```
SELECT * FROM employee  
WHERE upper(fname) = 'FRANKLIN';
```

```
SELECT essn FROM dependent WHERE  
age(d.bdate) > interval '18 years');
```



# BETWEEN and LIKE in SQL

- **BETWEEN, LIKE** are used in predicate:
  - **SELECT ... WHERE A BETWEEN 10 TO 20;**
  - **SELECT ... WHERE A1 LIKE '%IX%' OR A2 LIKE 'ABC%' OR A3 LIKE '%XYZ';**
  - **SELECT ... WHERE A1 LIKE '\_X\_%';**
- Also: **NOT BETWEEN** and **NOT LIKE**.



# Regular Expression Matching in PostgreSQL

- PostgreSQL also allows regular expression matching in string match using **IS SIMILAR TO <reg-ex>**





# Issues with Null Values



# Issues with Null Values

- An attribute having NULL could mean either of following-
  - Value is unknown or not available right now
  - Value is not application for the tuple: a employee not having supervisor will have null in this attribute
- Arithmetic expressions (+,-,\*,/) involving null values result null value for result



# Null Values and Comparisons

- There are more issues with Null value when it appears in a attribute used in WHERE clause of SQL-SELECT
- What should (  $r.a < 10$  ) where r.a happens to be null
- When we compare a NULL value with another value including NULL, result is UNKNOWN.
- In relational operations UNKNOWN is another truth-value like TRUE and FALSE



# Truth values for UNKNOWN

- **NOT**
  - NOT UNKNOWN -> UNKNOWN
- **AND**
  - TRUE AND UNKNOWN -> UNKNOWN
  - FALSE AND UNKNOWN -> FALSE
  - UNKNOWN AND UNKNOWN -> UNKNOWN
- **OR**
  - TRUE OR UNKNOWN -> TRUE
  - FALSE OR UNKNOWN -> UNKNOWN
  - UNKNOWN OR UNKNOWN -> UNKNOWN



# Null Values and Comparisons

- While evaluating WHERE clause tuples with UNKNOWN or FALSE truth values are not included in result
- Following query will not include any tuple where either of value is NULL irrespective of value in other attribute

```
SELECT * FROM EMPLOYEE WHERE  
bdate < DATE '2001-01-01' AND salary > 30000
```

- Following query will not include a tuple only when both are NULL, if one of attribute meets the condition then it will get included in result

```
SELECT * FROM EMPLOYEE WHERE  
bdate < DATE '2001-01-01' OR salary > 30000
```



# Null Values and Comparisons

## – IS NULL

- Following will not give desired result. Why? -

```
SELECT * FROM employee  
WHERE superssn = NULL;
```

- This is so because Null = Null is also UNKNOWN. For checking an attribute for having NULL value, SQL provides IS NULL (and IS NOT NULL)
- We write as following for such situations –  

```
SELECT * FROM employee  
WHERE superssn IS NULL ;
```



# Bags and Relational Operations



# Relational Operations and multiset (or bag)

- By Definition, relations are set; but implementations may permit duplicate tuples and such relations are called *bags*
- Normally stored relations (base) relations should still be sets, because most relations have PK
- However SQL SELECT results are often bags, possibly because duplicate removal is expensive.
- To get *set* you use DISTINCT keyword





# SQL and Multiset (or Bag)

- SET operations, that are UNION, INTERSECT, and EXCEPT in SQL yield their result as SET, that means duplicates are removed
- SQL however provides options by which you can have bag results by adding ALL keyword to operation name, i.e. UNION ALL, EXCEPT ALL or so.
- Let us see an example-



# UNION/INTERSECT/EXCEPT ALL in SQL

- Compare result of following queries:

```
SELECT superssn FROM employee; --Q1
```

```
SELECT mgrssn FROM department; --Q2
```

```
SELECT superssn FROM employee  
UNION
```

```
SELECT mgrssn FROM department; --Q3
```

```
SELECT superssn FROM employee  
UNION ALL
```

```
SELECT mgrssn FROM department; --Q4
```



R	superssn
	102
	101
	101
	102
	101
	(Null)
	108
	108

S	mgrssn
	101
	102
	108

R UNION S	superssn
	101
	102
	108
	(Null)

R UNION ALL S	superssn
	102
	101
	101
	102
	101
	(Null)
	108
	108
	101
	102
	108



R	superssn
	102
	101
	101
	102
	101
	(Null)
	108
	108

S	mgrssn
	101
	102
	108

R EXCEPT S	superssn
	(Null)

R EXCEPT ALL S	superssn
	101
	101
	102
	108
	(Null)



# UNION / INTERSECT / EXCEPT ALL in SQL

- UNION ALL
  - count of an element  $e$  in result is sum of count in  $R$  and  $S$
- INTERSECT ALL
  - $\min(\text{count-}r, \text{count-}s)$  of an element in  $R$  and  $S$ , is taken as result
- EXCEPT ALL:
  - Every occurrence of an element  $e$  in  $S$  decreases its count in  $R$  by one.



# Sub-queries in SQL



# Subquery in SQL

- A Query that is part of another query is *subquery*. A subquery may also have subquery, and so forth upto any level
- A subquery in SQL is written as a *query expression* enclosed in parentheses, and is in following form-  
**" ( SELECT ... FROM ... ) "**  
as a part of some existing query
- Result of sub-query is again a relation;



# Subquery in FROM clause

- FROM clause of SQL SELECT can have a sub-query, as following-

```
SELECT e.ssn, fname, dno, dname
FROM employee AS e NATURAL JOIN (SELECT
mgrssn AS ssn, dno, dname FROM
department) AS dept;
```

- In queries like above naming to relation returned by subquery is required, even if the relation name is not used; and that is the only relation as in query below-

```
SELECT * FROM (SELECT mgrssn AS ssn,
dno, dname FROM department) AS dept;
```





# Subquery in WHERE clause

- We have seen sub-query in **IN**, as  
**WHERE . . . IN ( SELECT . . . )**
- For example,  
**SELECT \* FROM employee WHERE ssn IN  
( SELECT essn FROM works\_on );**



# Subquery in WHERE clause

- When used in where clause and a sub-query returns a single column, single tuple relation, it can be interpreted as single value as following
- **SELECT pname FROM project  
WHERE dno = (SELECT dno FROM department  
WHERE dname = 'Research' );**
- Note: resultant relation of sub-query is getting compared with a attribute value – appropriate conversion takes place;
  - Has a underlying assumption that sub-query return a single tuple



# Execution of Subquery

- **SUB-Query may not execut for every tuple of outer query**
- Consider the query below-

```
SELECT pname FROM project
WHERE dno = (SELECT dno FROM department
WHERE dname = 'Research' );
```

- It is to typically executed as following: Execute inner query, and let us say returns **5**, and then places it in outer query, and the query to be executed becomes following -  

```
SELECT pname FROM project
WHERE dno = 5;
```



# Execution of Subquery

- **SUB-Query may not execute for every tuple of outer query**
- Consider another query-

```
SELECT * FROM student WHERE  
progid IN (SELECT pid FROM program  
           WHERE did = 'EE' );
```

- Typically, after execution of inner query, outer query may be translated to:

```
SELECT * FROM student WHERE  
progid IN (BEC, BEE);
```



# Correlated Sub-Queries



# Correlated Sub-Queries

- When inner query makes a reference to tuple of outer query then it is correlated sub-query. Consider following query -

- List employees, whose salary is more than department average:

```
SELECT ssn, fname FROM employee as e
WHERE salary > (SELECT AVG(salary) FROM
employee WHERE dno = e.dno)
```



# Execution of Correlated Sub-Queries

- Consider same query

```
SELECT ssn, fname FROM employee as e
WHERE salary > (SELECT AVG(salary)
FROM employee WHERE dno = e.dno)
```

- Logically, it is as following: For each tuple of outer query, execute inner query.
- Note that it can not be executed once for all tuples of outer query, as the case be with un-related inner query, and we have to execute **SUB-Query for every tuple of outer query**
- This is identified problem with correlated sub-queries.



## Correlated Sub-Queries could be expensive to execute – therefore should be avoided

- Correlated queries are expensive to execute, and can be avoided; for example the previous example
- **SELECT ssn, fname FROM employee as e  
WHERE salary > (SELECT AVG(salary)  
FROM employee WHERE dno = e.dno)**
- can be re-written as-  
**SELECT ssn, fname, salary FROM employee as e  
NATURAL JOIN (SELECT dno, AVG(salary) as  
avg\_sal FROM employee GROUP BY dno) as av  
WHERE salary > av.avg\_sal;**





# more Correlated Sub-queries

- List down employees having salary greater than their immediate supervisors.

```
select * from employee as e1 where e1.salary >
(select salary from employee as e2 where e2.ssn =
e1.superssn);
```

- Select employees having dependents older than 18 years:  

```
SELECT * FROM employee AS e WHERE ssn IN (SELECT
essn FROM dependent AS d WHERE essn = e.ssn AND
age(d.bdate) > interval '18 years');
```
- Attempt re-writting them without correlated query.



# EXISTS and NOT EXISTS in SQL

- Checks for emptiness of a relation and returns true or false.
- **EXISTS (r)** can be interpreted as “is there some tuple exists in relation r”
- **EXISTS (r)** returning true says that argument relation r is not empty
- Similarly, **NOT EXISTS (r)** returning true says that argument relation r empty



# Example EXISTS

- `SELECT ssn FROM employee AS e WHERE EXISTS  
(SELECT * FROM dependent AS d WHERE d.essn  
= e.ssn AND age(d.bdate) > interval '18  
years' );`



# SQL- EXISTS and IN

- While they might appear to be serving similar purposes, semantically are different.
- Both appear as part of predicate in WHERE clause of SELECT
- IN:
  - Syntax: **x IN ( r )**
  - Meaning: checks existence of tuple x in relation r, if found returns true, other wise false. Normally x is a scalar value and r is a single column relation.
- EXISTS:
  - Syntax: **EXISTS ( r )**
  - Meaning: checks if r is a non empty relation. Returns true if the relation has at least one tuple, otherwise false.
- In both above cases **r** is a *relational expression* resulting a relation.



# Compare a values with a bag of values (SQL)

- For example consider following two queries [Find out employee who have salary greater some or all employees of dno = 4]

SELECT ssn, fname FROM employee WHERE salary

> **SOME** (SELECT salary FROM employee WHERE dno = 4);

SELECT ssn, fname FROM employee WHERE salary

> **ALL** (SELECT salary FROM employee WHERE dno = 4);



# Compare a values with a bag of values (SQL)

- Note the equivalences:

SELECT ssn, fname FROM employee WHERE salary

> **SOME** (SELECT salary FROM employee WHERE dno = 4); and

SELECT ssn, fname FROM employee WHERE salary

> (SELECT **min**(salary) FROM employee WHERE dno = 4);

SELECT ssn, fname FROM employee WHERE salary

> **ALL** (SELECT salary FROM employee WHERE dno = 4); and

SELECT ssn, fname FROM employee WHERE salary

> (SELECT **max**(salary) FROM employee WHERE dno = 4);



# Compare a values with a bag of values (SQL)

- Comparative operators could be, one of following-  
**>SOME, >=SOME, <=SOME, <SOME, =SOME, <>SOME**  
**>ALL, >=ALL, <=ALL, <ALL, =ALL, <>ALL**
- Note: it can be easily proved that  
**=SOME** is identical to **IN**, and  
**<>SOME** is not identical to **NOT IN**  
**= ALL** is not identical to **IN**  
**<> ALL** (mean = NONE) and is same as **NOT IN**, and  
Earlier versions of SQL used **ANY** for **SOME**; today both keywords are used as synonymous.



# Sub-queries in Update statements

- `UPDATE employee`  
    `SET salary = salary * 1.1`  
    `WHERE ssn IN ( ... );`
- `DELETE employee`  
    `WHERE ssn = ( ... );`